

Cuckoo Hashing – Worst case $O(1)$ Lookup!

Background :

There are three basic operations that must be supported by a **hash table** (or a dictionary):

- **Lookup(key):** return **true** if key is there in the table, else **false**
- **Insert(key):** adds the item 'key' to the table if not already present
- **Delete(key):** removes 'key' from the table

Collisions are very likely even if we have a big table to store keys. Using the results from the **birthday paradox**: with only 23 persons, the probability that two people share the same birth date is 50%! There are 3 general strategies towards resolving hash collisions:

- **Closed addressing or Chaining:** store colliding elements in an auxiliary data structure like a linked list or a binary search tree.
- **Open addressing:** allow elements to overflow out of their target bucket and into other spaces.

Although above solutions provide expected lookup cost as $O(1)$, the expected worst-case cost of a lookup in Open Addressing (with linear probing) is $\Omega(\log n)$ and $\Theta(\log n / \log \log n)$ in simple chaining (Source : [Stanford Lecture Notes](#)). To close the gap of expected time and worst case expected time, two ideas are used:

- **Multiple-choice hashing:** Give each element multiple choices for positions where it can reside in the hash table
- **Relocation hashing:** Allow elements in the hash table to move after being placed

Cuckoo Hashing :

Cuckoo hashing applies the idea of multiple-choice and relocation together and guarantees $O(1)$ worst case lookup time!

- **Multiple-choice:** We give a key **two choices** $h_1(\text{key})$ and $h_2(\text{key})$ for residing.
- **Relocation:** It may happen that $h_1(\text{key})$ and $h_2(\text{key})$ are preoccupied. This is resolved by imitating the Cuckoo bird: *it pushes the other eggs or young out of the nest when it hatches*. Analogously, inserting a new key into a cuckoo hashing table may push an older key to a different location. This leaves us with the problem of re-placing the older key.
 - If alternate position of older key is vacant, there is no problem.
 - Otherwise, older key displaces another key. This continues until the procedure finds a vacant position, or enters a cycle. In case of cycle, new hash functions are chosen and the whole data structure is 'rehashed'. Multiple rehashes might be necessary before Cuckoo succeeds.

Insertion is expected $O(1)$ (amortized) with high probability, even considering the possibility rehashing, as long as the number of keys is kept below half of the capacity of the hash table, i.e., the load factor is below 50%.

Deletion is $O(1)$ worst-case as it requires inspection of just two locations in the hash table.

Illustration :

Input:

```
{20, 50, 53, 75, 100, 67, 105, 3, 36, 39}
```

Hash Functions:

```
h1(key) = key%11  
h2(key) = (key/11)%11
```

	20	50	53	75	100	67	105	3	36	39
$h_1(\text{key})$	9	6	9	9	1	1	6	3	3	6
$h_2(\text{key})$	1	4	4	6	9	6	9	0	3	3

Let's start with inserting **20** at its possible position in the first table determined by $h_1(20)$:

table[1]	-	-	-	-	-	-	-	-	-	20	-
table[2]	-	-	-	-	-	-	-	-	-	-	-

Next: **50**

table[1]	-	-	-	-	-	-	50	-	-	20	-
table[2]	-	-	-	-	-	-	-	-	-	-	-

Next: **53**. $h_1(53) = 9$. But 20 is already there at 9. We place 53 in table 1 & 20 in table 2 at $h_2(20)$

table[1]	-	-	-	-	-	-	50	-	-	53	-
table[2]	-	20	-	-	-	-	-	-	-	-	-

Next: **75**. $h_1(75) = 9$. But 53 is already there at 9. We place 75 in table 1 & 53 in table 2 at $h_2(53)$

table[1]	-	-	-	-	-	-	50	-	-	75	-
table[2]	-	20	-	-	53	-	-	-	-	-	-

Next: **100**. $h_1(100) = 1$.

table[1]	-	100	-	-	-	-	50	-	-	75	-
table[2]	-	20	-	-	53	-	-	-	-	-	-

Next: **67**. $h_1(67) = 1$. But 100 is already there at 1. We place 67 in table 1 & 100 in table 2

table[1]	-	67	-	-	-	-	50	-	-	75	-
table[2]	-	20	-	-	53	-	-	-	-	100	-

Next: **105**. $h_1(105) = 6$. But 50 is already there at 6. We place 105 in table 1 & 50 in table 2 at $h_2(50) = 4$. Now 53 has been displaced. $h_1(53) = 9$. 75 displaced: $h_2(75) = 6$.

table[1]	-	67	-	-	-	-	105	-	-	53	-
table[2]	-	20	-	-	50	-	75	-	-	100	-

Next: **3**. $h_1(3) = 3$.

table[1]	-	67	-	3	-	-	105	-	-	53	-
table[2]	-	20	-	-	50	-	75	-	-	100	-

Next: **36**. $h_1(36) = 3$. $h_2(3) = 0$.

table[1]	-	67	-	36	-	-	105	-	-	53	-
table[2]	3	20	-	-	50	-	75	-	-	100	-

Next: **39**. $h_1(39) = 6$. $h_2(105) = 9$. $h_1(100) = 1$. $h_2(67) = 6$. $h_1(75) = 9$. $h_2(53) = 4$. $h_1(50) = 6$. $h_2(39) = 3$.

Here, the new key 39 is displaced later in the recursive calls to place 105 which it displaced.

table[1]	-	100	-	36	-	-	50	-	-	75	-
table[2]	3	20	-	39	53	-	67	-	-	105	-

Implementation :

Below is C/C++ implementation of Cuckoo hashing

```
// C++ program to demonstrate working of Cuckoo
// hashing.
#include<bits/stdc++.h>

// upper bound on number of elements in our set
#define MAXN 11

// choices for position
#define ver 2

// Auxiliary space bounded by a small multiple
// of MAXN, minimizing wastage
int hashtable[ver][MAXN];

// Array to store possible positions for a key
int pos[ver];

/* function to fill hash table with dummy value
 * dummy value: INT_MIN
 * number of hashtables: ver */
void initTable()
{
    for (int j=0; j<MAXN; j++)
        for (int i=0; i<ver; i++)
            hashtable[i][j] = INT_MIN;
}

/* return hashed value for a key
 * function: ID of hash function according to which
 * key has to hashed
 * key: item to be hashed */
int hash(int function, int key)
{
    switch (function)
    {
        case 1: return key%MAXN;
        case 2: return (key/MAXN)%MAXN;
    }
}

/* function to place a key in one of its possible positions
 * tableID: table in which key has to be placed, also equal
 * to function according to which key must be hashed
 * cnt: number of times function has already been called
 * in order to place the first input key
 * n: maximum number of times function can be recursively
 * called before stopping and declaring presence of cycle */
void place(int key, int tableID, int cnt, int n)
{
    if (hashtable[tableID][key] == INT_MIN)
    {
        hashtable[tableID][key] = key;
        return;
    }
    if (cnt == n)
    {
        cout << "Cycle detected\n";
        return;
    }
    int pos1 = hash(tableID, key);
    int pos2 = hash(1 - tableID, key);
    if (pos1 == pos2)
    {
        cout << "Cycle detected\n";
        return;
    }
    if (pos1 == tableID)
        place(key, pos2, cnt + 1, n);
    else
        place(key, pos1, cnt + 1, n);
}
```

```

/* if function has been recursively called max number
of times, stop and declare cycle. Rehash. */
if (cnt==n)
{
    printf("%d unpositioned\n", key);
    printf("Cycle present. REHASH.\n");
    return;
}

/* calculate and store possible positions for the key.
* check if key already present at any of the positions.
If YES, return. */
for (int i=0; i<ver; i++)
{
    pos[i] = hash(i+1, key);
    if (hashtable[i][pos[i]] == key)
        return;
}

/* check if another key is already present at the
position for the new key in the table
* If YES: place the new key in its position
* and place the older key in an alternate position
for it in the next table */
if (hashtable[tableID][pos[tableID]]!=INT_MIN)
{
    int dis = hashtable[tableID][pos[tableID]];
    hashtable[tableID][pos[tableID]] = key;
    place(dis, (tableID+1)%ver, cnt+1, n);
}
else //else: place the new key in its position
    hashtable[tableID][pos[tableID]] = key;
}

/* function to print hash table contents */
void printTable()
{
    printf("Final hash tables:\n");

    for (int i=0; i<ver; i++, printf("\n"))
        for (int j=0; j<MAXN; j++)
            (hashtable[i][j]==INT_MIN)? printf("- "):
            printf("%d ", hashtable[i][j]);

    printf("\n");
}

/* function for Cuckoo-hashing keys
* keys[]: input array of keys
* n: size of input array */
void cuckoo(int keys[], int n)
{
    // initialize hash tables to a dummy value (INT-MIN)
    // indicating empty position
    initTable();

    // start with placing every key at its position in
    // the first hash table according to first hash
    // function
    for (int i=0, cnt=0; i<n; i++, cnt=0)
        place(keys[i], 0, cnt, n);

    //print the final hash tables
    printTable();
}

/* driver function */
int main()
{
    /* following array doesn't have any cycles and
    hence all keys will be inserted without any
    rehashing */

```

```

int keys_1[] = {20, 50, 53, 75, 100, 67, 105,
                3, 36, 39};

int n = sizeof(keys_1)/sizeof(int);

cuckoo(keys_1, n);

/* following array has a cycle and hence we will
   have to rehash to position every key */
int keys_2[] = {20, 50, 53, 75, 100, 67, 105,
                3, 36, 39, 6};

int m = sizeof(keys_2)/sizeof(int);

cuckoo(keys_2, m);

return 0;
}

```

Output:

```

Final hash tables:
- 100 - 36 - - 50 - - 75 -
3 20 - 39 53 - 67 - - 105 -

105 unpositioned
Cycle present. REHASH.
Final hash tables:
- 67 - 3 - - 39 - - 53 -
6 20 - 36 50 - 75 - - 100 -

```

Generalizations of cuckoo hashing that use more than 2 alternative hash functions can be expected to utilize a larger part of the capacity of the hash table efficiently while sacrificing some lookup and insertion speed. Example: if we use 3 hash functions, it's safe to load 91% and still be operating within expected bounds (Source : [Wiki](#))